

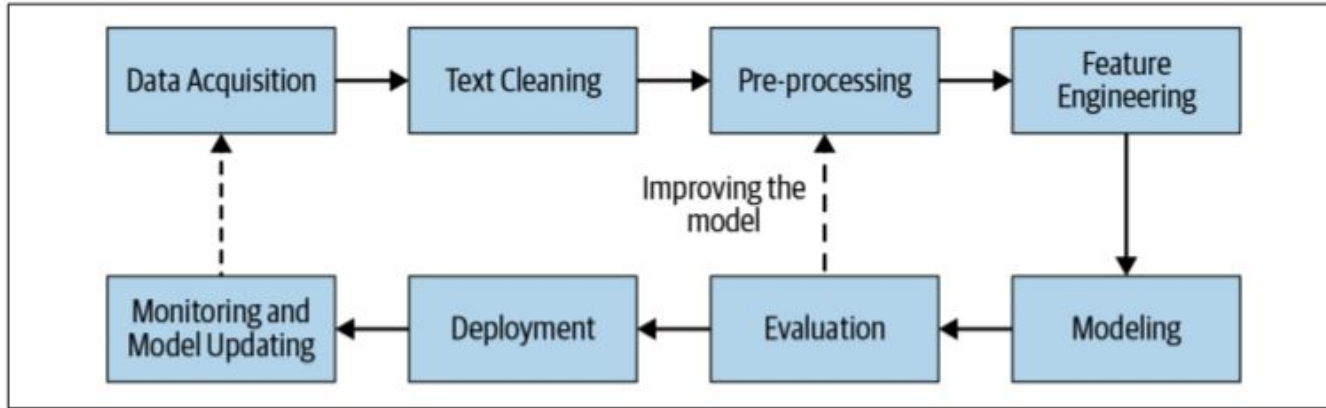
NLP Applications II

NLP Pipeline

Introduction

- When we build an application, we normally walk through the requirements and **break the problem down into several sub-problems**.
- The step-by-step processing of text is known as a **pipeline**.
 - It is the series of steps involved in building any NLP model.
 - These steps are **common in every NLP project**
- We will learn about these steps and how **they play important roles** in solving NLP problems.

Generic Pipeline



Key steps in NLP

1. Data acquisition
2. Text cleaning
3. Pre-processing
4. Feature engineering
5. Modeling
6. Evaluation
7. Deployment
8. Monitoring and model updating

Figure 2-1. Generic NLP pipeline

Source: <http://www.practicalnlp.ai/>

1. Data Acquisition

- In an ideal setting we would have a dataset with **thousand or millions of data points**.
- **Realistic setting** (most of the industrial projects)
 - Lack of data becomes the bottleneck of many projects.
 - Important to know about different ways to gather data.
- **Common scenario**: Many companies have a large amount of PDF documents, but no document is manually annotated.
- So how can we get useful annotated data?

Strategies for data acquisition

- Public datasets
 - Find a similar task and domain (see <https://www.kaggle.com/>, <https://datasets.quantumstat.com/>, <https://datasetsearch.research.google.com/>)
 - Build model and evaluate on your problem
- Scrape data from Internet
 - Extract data e.g discussion forum
 - Label/annotate data
- Product intervention
 - Get data from your company
 - E.g Netflix, Facebook, Microsoft, Google, Amazon...
- Data augmentation
 - Start from small dataset, create automatically more data

Data augmentation

Replacement strategies

- Synonym replacement
- TF-IDF based word replacement
- (Named) entity replacement

Adding noise

- Bigram flipping
- Character level noise
- Random insertion

Einstein was one of the **outstanding** figures of the 20th century"



Bohr was one of the **great** figures of the 20th century"

I am **going to** the supermarket




I am **to going** the supermarket

EDA: data augmentation

EDA: Easy Data Augmentation Techniques for Boosting Performance on Text Classification Tasks. Jason Wei, Kai Zou. EMNLP-IJCNLP. 2019

- paper: <https://arxiv.org/abs/1901.11196>
- code: https://github.com/jasonwei20/eda_nlp



- **Synonym Replacement (SR):** Randomly choose n words from the sentence. Replace each of these words with one of its synonyms chosen at random.
- **Random Insertion (RI):** Find a random synonym of a random word in the sentence. Insert that synonym into a random position in the sentence. Do this n times.
- **Random Swap (RS):** Randomly choose two words in the sentence and swap their positions. Do this n times.
- **Random Deletion (RD):** For each word in the sentence, randomly remove it with probability p .

Data augmentation: Backtranslation

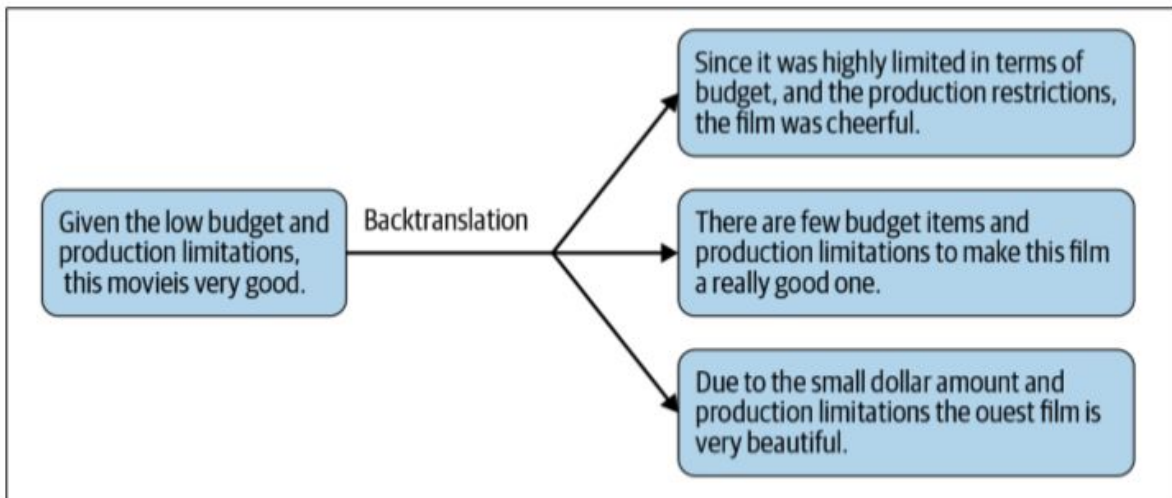


Figure 2-2. Back translation

- Use MT to generate paraphrases
 - MT: en→ es; MT: es→ en
- Easy to use MT implementations
 - <https://github.com/pytorch/fairseq>
- Gain more variety of augmented examples

2. Text extraction and cleaning

- Remove all non-textual information:
 - Mark-up., metadata
 - OCR
- HTML, PDF to text
- Do not require any NLP-specific technique
- **Most of the data is not just raw text!**




Figure 2-1. Generic NLP pipeline

A

B

```

<div style="text-align: justify;">
  The book will be around 350 pages. It will be accompanied by a code
  repository containing several Jupyter notebooks for all the chapters to
  give a walk-through and explain the code in detail. The code base is in
  Python and various machine learning and natural language processing
  libraries. The book assumes that the readers have a good grasp of
  programming but no theoretical and practical knowledge of NLP.
</div>

<p><br>
</p>

<a
  href="https://www.oreilly.com/library/view/practical-natural-language/97
  81492054047/">

</a>

<h1 style="color: #e74c3c;">Commonly Asked Questions</h1>
<ul>
<li>Can I contribute to the book?
<p>The book is accompanied by open source Jupyter notebooks and demo
  applications. If you are a great ML or front-end engineer looking to
  build something meaningful you can apply by filling <a
  href="https://goo.gl/forms/d0fNcw251IX26ajk1">this form</a>. Also refer
  to the next question.
            
```

C

Figure 2-3. (a) PDF Invoice [13] (b) HTML texts (c) text embedded in an image [14]

3. Preprocessing


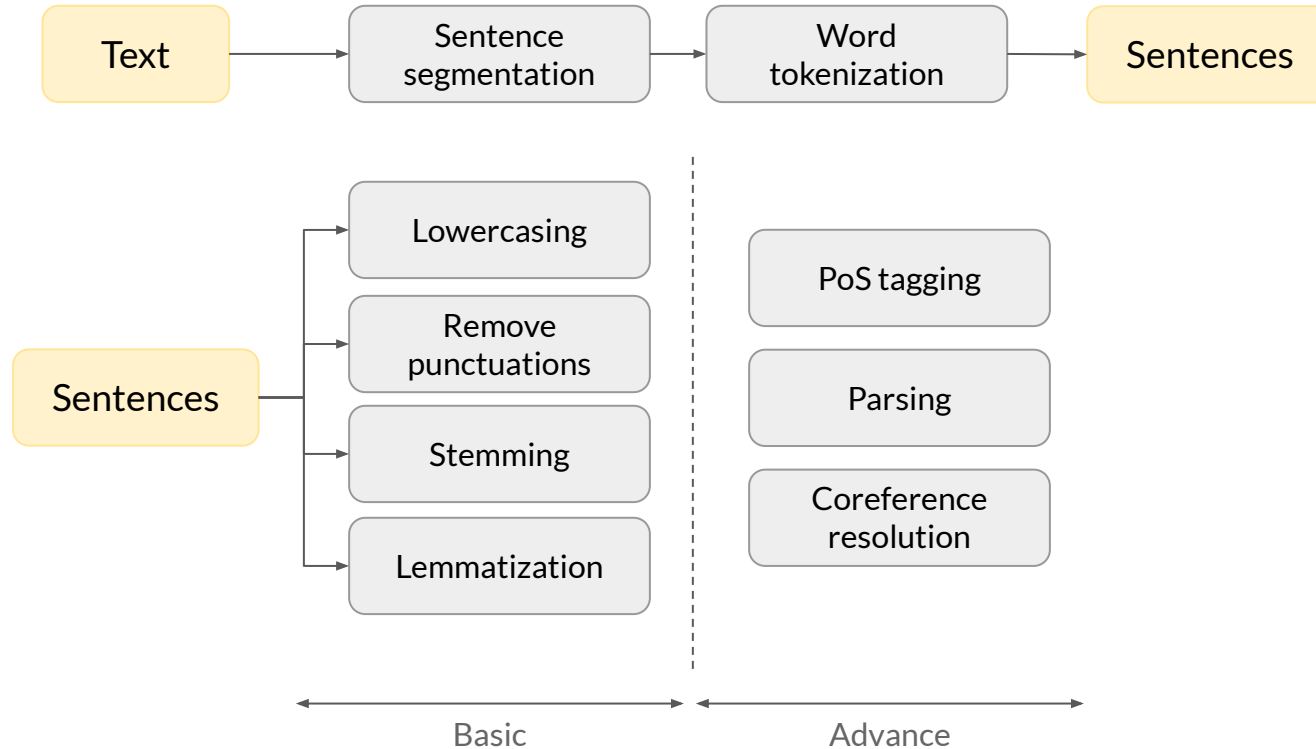


Figure 2-1. Generic NLP pipeline



Preprocessing tools

- Natural Language Toolkit (NLTK): <http://www.nltk.org/>
- SpaCy: <https://spacy.io/>
- StanfordNLP: <https://stanfordnlp.github.io/stanfordnlp/>
- Trankit: <http://nlp.uoregon.edu/trankit>
- **Similar way of use:**
 - Load text processor (tokenizer, tagger, NER)
 - Process document or text chunk
 - Extract the linguistic information you need
- **Many pretrained models and languages!**

Edit the code & try spaCy

spaCy v3.0 · Python 3 · via Binder

```
# pip install -U spacy
# python -m spacy download en_core_web_sm
import spacy

# Load English tokenizer, tagger, parser and NER
nlp = spacy.load("en_core_web_sm")

# Process whole documents
text = ("When Sebastian Thrun started working on self-driving cars at "
        "Google in 2007, few people outside of the company took him "
        "seriously. "I can tell you very senior CEOs of major American "
        "car companies would shake my hand and turn away because I wasn't "
        "worth talking to," said Thrun, in an interview with Recode earlier "
        "this week.")

doc = nlp(text)

# Analyze syntax
print("Noun phrases:", [chunk.text for chunk in doc.noun_chunks])
print("Verbs:", [token.lemma_ for token in doc if token.pos_ == "VERB"])

# Find named entities, phrases and concepts
for entity in doc.ents:
    print(entity.text, entity.label_)
```

RUN

```
>>> import stanfordnlp
>>> stanfordnlp.download('en') # This downloads the English models for the neural pipeline
>>> nlp = stanfordnlp.Pipeline() # This sets up a default neural pipeline in English
>>> doc = nlp("Barack Obama was born in Hawaii. He was elected president in 2008.")
>>> doc.sentences[0].print_dependencies()
```

```
1 from trankit import Pipeline
2 # initialize a pipeline on English
3 p = Pipeline(lang='english', gpu=True, cache_dir='./cache')
4
5 doc = '''Michael helped shoot the majority of my firm's website
6 and we could not have been happier.'''
7
8 # perform all tasks on the input
9 all = p(doc)
10
11 sents = p.ssplit(doc) # sentence segmentation
12 tokens = p.tokenize(doc) # tokenization
13 posdeps = p.posdep(doc) # upos, xpos, ufeats, dependency parsing
14 ners = p.ner(doc) # ner tagging
15 lemmas = p.lemmatize(doc) # Lemmatization
```

4. Feature engineering




Figure 2-1. Generic NLP pipeline

- Set of methods that feed pre-processed text into ML algorithms.
- Capture the characteristics of the text into a numeric vector that can be understood by the ML algorithms.
- Also known as *feature extraction* or *text representation*.


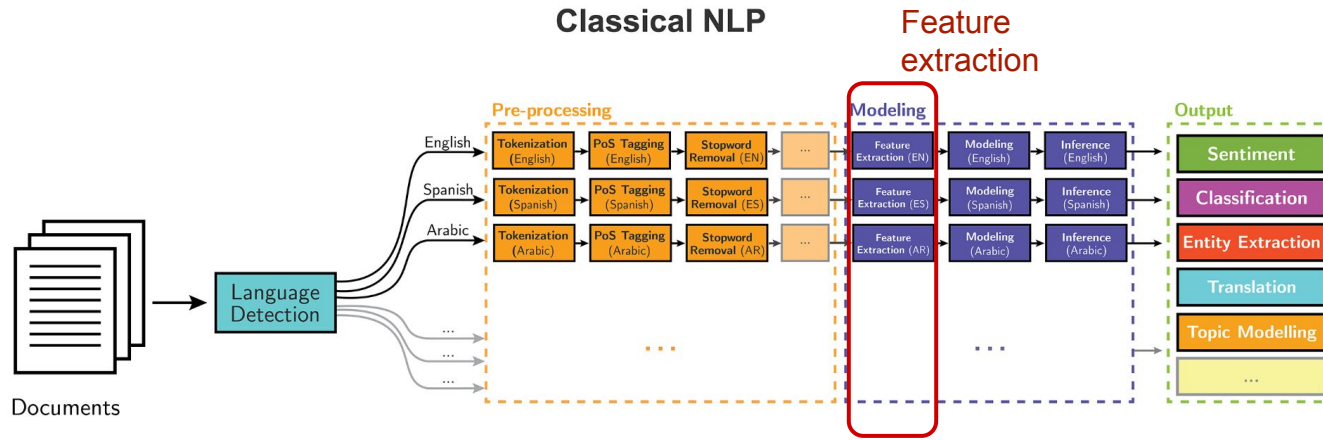
Two different approaches:

1. Classical NLP with traditional ML pipeline

- Hand crafted features: e.g. `number_positive_words(sentence)`
- Domain knowledge
- Statistical measure to usefulness of features (interpretable)
- Noisy/unrelated information

2. DL pipeline

- Capable of learning *features* (hidden) inline with task requirements
- Difficult to interpret



5. Modeling

Simple Heuristics

- Start building systems by encoding this knowledge in the form of rules/heuristics
- Spam detection. Use blacklist to learn words associated with spam emails
- Regular expression to extract phone numbers, names, etc.
- <https://spacy.io/usage/rule-based-matching> (rule-based matching)
- <https://explosion.ai/demos/matcher>

```
{{([ { ner:PERSON } ] ) /was/ /born/ /on/ ( [ { ner:DATE } ] ) }  
=>"DATE_OF_BIRTH" }
```

Use heuristics as features

- Combination of many heuristic can be fuzzy
- Use heuristics as features (e.g. number of words from the blacklist)




Figure 2-1. Generic NLP pipeline

Modeling

AutoML

- provides automatic methods and processes to apply Machine Learning
- data pre-processing, feature engineering, algorithm selection, hyperparameter optimization
- Make ML available for non-Machine Learning experts

Some examples:

- [AutoWEKA](#): simultaneous selection of a machine learning algorithm and its hyperparameters.
- [Auto-sklearn](#): similar framework for [scikit-learn](#)
- [H2O AutoML](#): provides automated model selection and ensembling for the [H2O machine learning and data analytics platform](#).
- [MLBoX](#): AutoML library with three components: preprocessing, optimisation and prediction.

Modeling

Ensemble and stacking

- Combining more than one prediction usually improve results
- **Model stacking:** Feed one model's output as input for another.
- **Model ensembling:** Pool predictions from multiple models and make final prediction.

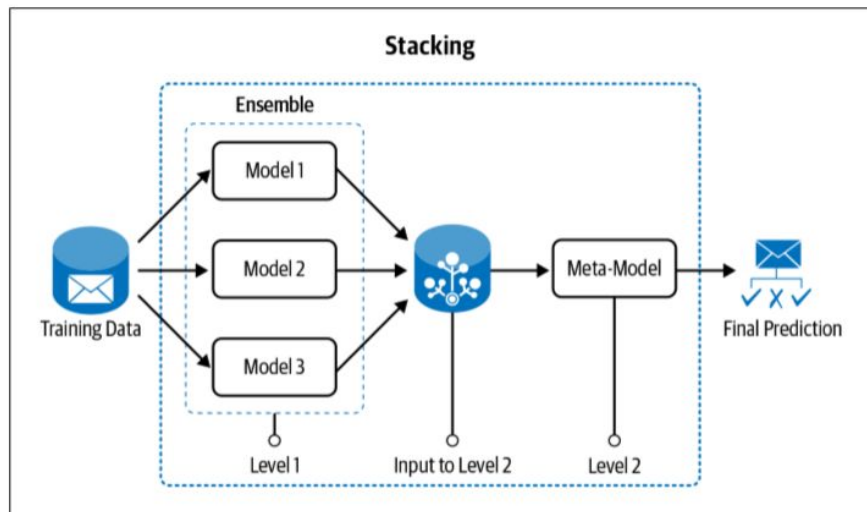


Figure 2-14. Model ensemble and stacking

Modeling

Better feature engineering

- Using many feature can be inconvenient
- Use **feature selection** to find a better model

Transfer learning

- Transfer preexisting knowledge from a big, well-trained model to a newer model at its initial phase
- Provides a better initialization, which helps in the downstream tasks, especially when the dataset for the downstream task is smaller.

Modeling

Large data volume

- Use DL or a richer set of features

Small data volume

- Start with rule-based or traditional ML.
- Use transfer learning if possible (need pre-trained LM)
- Use prompt learning if possible (need pre-trained model)

6. Evaluation

- Goodness of model can have multiple meanings.
 - The most common interpretation is the measure of the model's performance on unseen data.
- Success of evaluation depends on **evaluation metric and process**.
- **Evaluation metric** depends on task and phase
 - Model building, deployment, and production phases.
- Two types of evaluations:
 - **Intrinsic**: Focus on intermediary objective
 - **Extrinsic**: Focus on final objective
- E.g. spam-classification system.
 - The intrinsic metric will be precision and recall.
 - The business metric (extrinsic) will be “the amount of time users spent on a spam email.”

Intrinsic evaluation

- Assume a test set with a ground-truth/labels (human annotated examples)
- Labels can be binary (text classification), one-to-two words (e.g. NER), or large texts (e.g MT)
- The output of the NLP model is compared against the corresponding label for that data point.
- Metrics are calculated based on the match (or mismatch) between the output and label.
 - **Text classification, NER, Relation Extraction:** Precision, Recall, Fscore, ROC,
 - **Text generation:** BLEU, METEOR, ROUGE
 - **Similarity:** Pearson correlation, MSE
 - **Ranking:** MAP, Recall@K

Extrinsic evaluation

- Evaluation of model performance on **final practical objective**
- Model with great intrinsic metric can fail achieving business objective
 - QA system makes great on SQUAD, but might fail answering large number of question in the production environment.
- More **expensive than intrinsic evaluation** (that's why we need intrinsic evaluation).
- Bad results in intrinsic imply bad results in extrinsic.
- Good results in intrinsic do not imply good performance in extrinsic.

Wrap up

- We saw different steps involved in developing an NLP pipeline
- Specific details for each step will depend on the task at hand and the purpose of our implementation
- Differences between traditional NLP pipeline and a DL-based NLP pipeline
- Other languages might be treated differently.